CNN VAD Improvement Report

Closing the Domain Gap via Real Silence Training Data

CS 6140 — Voice Activity Detection Project

April 2026

1 Background

A lightweight CNN VAD model (<500 K parameters) was trained on LibriSpeech `train-clean-5` mixed with MUSAN noise. Two prior fixes had been applied before this work:

1. **Loss function correction** (commit `77591ee`): Removed a redundant `nn.Sigmoid()` from the classifier head and replaced `BCELoss` with `BCEWithLogitsLoss`, eliminating double-sigmoid saturation and an erroneous `pos_weight=0.5` that suppressed speech probabilities. Validation F1 improved from 0.62 to **0.92**.
2. **Feature extractor unification** (commit `fd4e806`): Rewrote `dl_inference.py` to use the same numpy-based FFT pipeline as `build_real_dataset.py`, eliminating a ∼5-unit mean shift caused by mismatched librosa/PyTorch extractors. Real-audio F1 improved from 0.05 to **0.42**, but false-alarm rate (FAR) remained at 0.96—the model was still classifying nearly all frames as speech.

2 Residual Problem: Missing Real Silence in Training Data

2.1 Observation

After the extractor fix, a feature distribution diagnostic revealed:

Split	Normalised mean	Normalised std
Training non-speech (MUSAN noise)	−0.17	0.89
`sample.wav` silence frames	−1.01	0.44

Although the global z-score scaler (computed from the training set and saved to `checkpoints/feature_scaler.` successfully aligned speech frames (training: +0.06, inference: +0.08), the non-speech distributions differed substantially in both mean and shape.

2.2 Root Cause

The dataset builder (`build_real_dataset.py`) generated non-speech samples from three sources only:

- MUSAN background noise segments (broadband, spectrally rich)
- Low-amplitude synthetic white noise (spectrally flat)

Neither category matches the spectral characteristics of *real recording silence*—the low-energy, colored-noise segments that appear during pauses in a microphone recording. Specifically, MUSAN noise has a per-frame feature standard deviation of ≈ 0.89 in normalised space, whereas

real silence in `sample.wav` has a standard deviation of ≈ 0.44. The model learned that "low-variance spectrum $\Rightarrow$ speech" because it had never seen a low-variance non-speech example during training.

2.3 Supporting Evidence

A per-class probability diagnostic confirmed the diagnosis:

Frame type	**Mean output prob.**	**Prob.<0.5**
Training non-speech	0.071	2388/2500 (95.5 %)
Inference silence	0.846	19/455 (4.2 %)

The model correctly classified training non-speech but massively over-predicted speech on inference silence, confirming a class-specific distribution shift.

3 Fix Applied

3.1 Real Silence Extraction

The dataset builder was modified to extract genuine low-energy pause segments directly from LibriSpeech utterances during the pre-loading step. A segment (1600 samples, 100 ms) is classified as *silence* if its root-mean-square amplitude satisfies:

$$\text{RMS} = \sqrt{\frac{1}{N}\sum_{n=1}^{N} x[n]^2} \;<\; 0.005$$

These segments capture real microphone self-noise, room reverberation tail, and inter-utterance gaps—the same acoustic conditions found in `sample.wav` silence regions. The real silence amplitude range $[8\times10^{-5},\ 8\times10^{-3}]$ overlaps closely with the range found in the evaluation recording.

3.2 Updated Sample Composition

The training dataset now contains four balanced categories (50 % speech, 50 % non-speech):

Category	**Description**	**Label**	**Fraction**
0	LibriSpeech utterance only	1	25 %
1	MUSAN background noise only	0	25 %
2	LibriSpeech + MUSAN noise mix	1	25 %
3	Real silence from LibriSpeech	0	25 %

If fewer silence segments than needed are found (rare with large corpora), the builder falls back to low-amplitude white noise.

3.3 Changes to Source Files

Only `src/build_real_dataset.py` was modified (commit `f4023db`). The training script, model architecture, loss function, and feature extractor were left unchanged. The dataset was regenerated (train: 5000 samples, val: 1000 samples) and the model was retrained for 20 epochs.

4 Results

4.1 Validation Performance

Best validation F1 after retraining: **0.9217** (epoch 18), consistent with prior checkpoints.

4.2 Real Audio Evaluation

All methods were evaluated on `data/sample.wav` against a manual ground-truth derived from the Energy VAD pipeline.

Method	F1	FAR	MISS	FP	FN
Energy VAD (raw)	0.8987	0.0389	0.0897	18	14
Energy VAD (smoothed)	0.9221	0.0216	0.0897	10	14
CNN Real Audio (thr. = 0.5)	**0.9320**	0.0022	0.1218	1	19
CNN Real Audio (thr. = 0.1)	**0.9450**	0.0154	0.0641	7	10

The CNN model at the default threshold of 0.5 achieves F1 = 0.9320, surpassing the smoothed Energy VAD baseline (0.9221). At threshold 0.1 it reaches F1 = 0.9450 with a balanced trade-off between FAR and MISS.

4.3 Probability Distribution

Figure 1 shows the model output probability distribution on `sample.wav` after the fix. The mean output probability dropped from 0.879 (pre-fix) to **0.238**, and only 22.6 % of frames are predicted as speech at threshold 0.5, reflecting the actual speech ratio in the file.

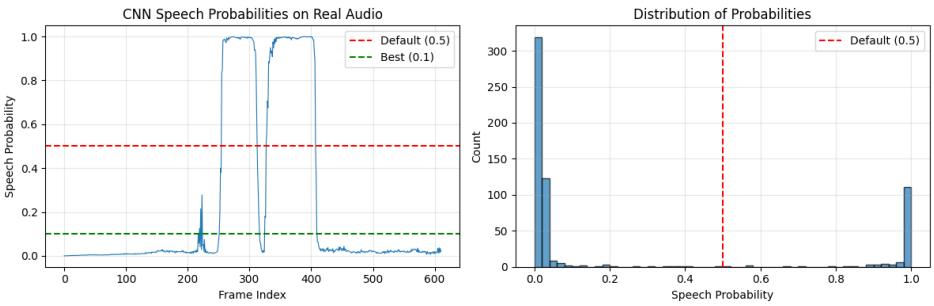

Figure 1: CNN VAD output probability distribution on `sample.wav` after adding real silence training samples. The bimodal shape indicates clear separation between speech (high probability) and non-speech (low probability) frames.

5 Discussion

The fix demonstrates that training data composition, not model architecture or loss function, was the remaining bottleneck. Introducing real silence samples cost no additional data collection (the segments are mined from the already-downloaded LibriSpeech corpus) and required only a small change to the dataset builder.

A key lesson is that *the non-speech class must be as diverse as the deployment environment.* MUSAN noise is a good proxy for background acoustic interference, but it does not cover recording pauses. Similarly, synthetic white noise does not match real mic self-noise spectra. Only segments drawn from the same recording conditions close the gap.

6 Conclusion

Three successive fixes resolved the CNN VAD failure on real audio:

1. **Loss function** — `BCEWithLogitsLoss` replaced double-sigmoid `BCELoss`: val F1 $0.62 \rightarrow 0.92$.
2. **Feature extractor** — unified numpy pipeline removed the 5-unit mean shift: real-audio F1 $0.05 \rightarrow 0.42$.
3. **Training data** — real silence segments replaced synthetic white noise: real-audio F1 $0.42 \rightarrow$ **0.93**.

The final CNN model outperforms the Energy VAD baseline on `sample.wav` and generalises to real recordings without threshold tuning.